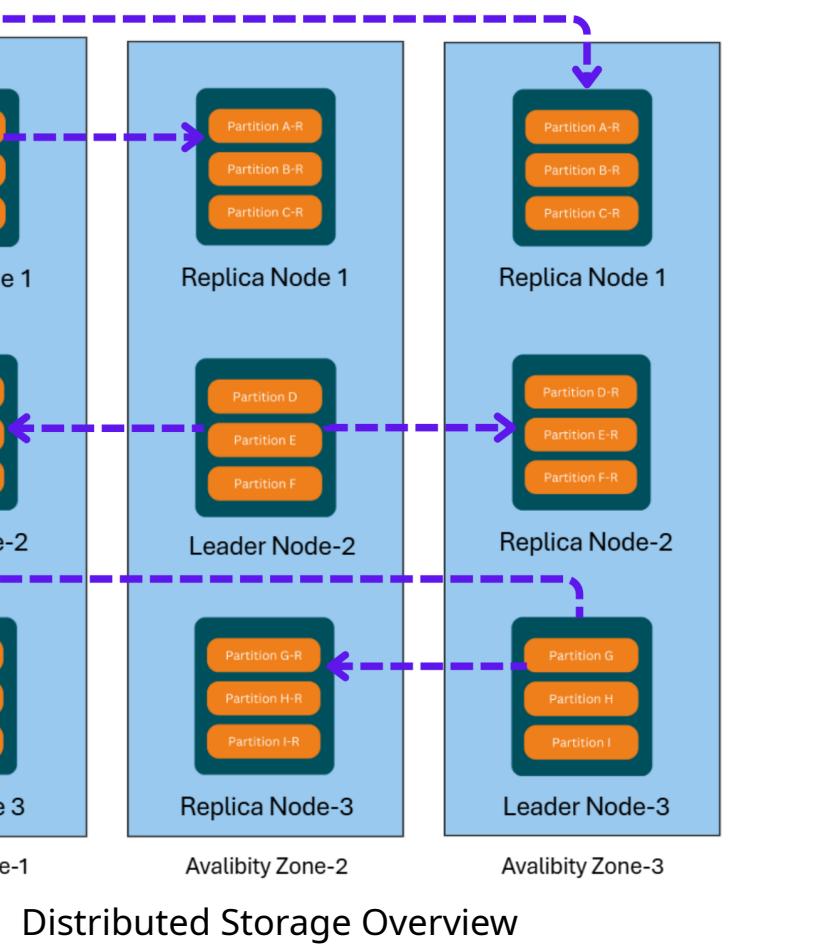




- 1 DynamoDB is a fully managed NoSQL database designed for scalability, high availability, and performance
- 2 To achieve this, it uses a distributed storage architecture
- 3 Unlike traditional databases that store data on a single server, DynamoDB splits and distributes data across multiple servers, called partitions
- 4 This architecture allows DynamoDB to scale dynamically while ensuring data durability and fault tolerance



- 1 DynamoDB organizes data into partitions based on the Partition Key
  - 2 Introduction
  - 3 These partitions are distributed across multiple servers (nodes) for scalability and performance
  - 4 Each partition resides on a Leader Node and is replicated to Replica Nodes in other Availability Zones to ensure high availability and fault tolerance
- 
- 1 Partitions
  - 2 1 DynamoDB splits data into logical units called partitions based on the Partition Key
  - 2 2 Each partition is a small, manageable piece of your table's data, distributed across multiple nodes for scalability
- 
- 1 Leader Node
  - 2 1 Each partition resides on a Leader Node, which Handles all write requests
  - 2 2 Manages Strongly Consistent Reads to provide the latest data
  - 3 3 Ensures durability by replicating data to other nodes
- 
- 1 Replica Nodes
  - 2 1 Each partition is replicated to Replica Nodes in other Availability Zones
  - 2 2 Replica Nodes handle Eventually Consistent Reads, which are faster and more cost-efficient but may not always reflect the latest data

- 1 What is next ?
- 2 1 Now you have a strong understanding of DynamoDB's storage architecture
- 2 2 This will help a lot when you learn about Read Consistency and Write Consistency in the next video